

model_quantization_xilinx

March 26, 2021

1 Model quantization

To perform model quantization, we use the Xilinx DNNDK tool (https://www.xilinx.com/support/documentation/sw_manuals/ai_inference/v1_6/ug1327-dnndk-user-guide.pdf).

```
[1]: import os
import sys
sys.path.append(os.path.dirname(os.getcwd()))
```

```
[4]: import os
import subprocess

import tensorflow as tf

from scripts.quantization import evaluate_graph, freeze_model
from scripts import prepare_data, artifacts_reporter, train_model,
    ↪ evaluate_model
from ml_intuition.data.utils import plot_training_curve, show_statistics
```

```
[5]: DEST_PATH = 'xilinx_model_compilation_results'
DATA_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    ↪ pavia.npy')
GT_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    ↪ pavia_gt.npy')
experiment_dest_path = os.path.join(DEST_PATH, 'experiment_0')
data_path = os.path.join(experiment_dest_path, 'data.h5')
os.makedirs(experiment_dest_path, exist_ok=True)
```

2 Prepare the data

To fit into the the pipeline, the data has to be preprocessed. It is achieved by the `prepare_data.main` function. It accepts a path to a `.npy` file with the original cube as well as the corresponding ground truth. In this example, we randomly extract 250 samples from each class (balanced scenario), use 10% of them as validation set, and extract only spectral information of a

pixel. The returned object is a dictionary with three keys: `train`, `test` and `val`. Each of them contains an additional dictionary with `data` and `labels` keys, holding corresponding `numpy.ndarray` objects with the data. For more details about the parameters, refer to the documentation of `prepare_data.main` function (located in `scripts/prepare_data`).

```
[6]: prepare_data.main(data_file_path=DATA_FILE_PATH,
                        ground_truth_path=GT_FILE_PATH,
                        output_path=data_path,
                        train_size=250,
                        val_size=0.1,
                        stratified=True,
                        background_label=0,
                        channels_idx=2,
                        neighborhood_size=None,
                        save_data=True,
                        seed=0)
```

3 Train the model

The function `train_model.train` executed the training procedure. Trained model will be stored under `experiment_dest_path` folder path.


```
[7]: train_model.train(model_name='model_2d',
                        kernel_size=5,
                        n_kernels=200,
                        n_layers=1,
                        dest_path=experiment_dest_path,
                        data=data_path,
                        sample_size=103,
                        n_classes=9,
                        lr=0.001,
                        batch_size=128,
                        epochs=200,
                        verbose=0,
                        shuffle=True,
                        patience=15,
                        noise=[],
                        noise_sets=[])
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 99, 1, 200)	1200
conv2d_1 (Conv2D)	(None, 32, 1, 200)	200200

conv2d_2 (Conv2D)	(None, 14, 1, 200)	200200
conv2d_3 (Conv2D)	(None, 5, 1, 200)	200200
flatten (Flatten)	(None, 1000)	0
dense (Dense)	(None, 200)	200200
dense_1 (Dense)	(None, 128)	25728
dense_2 (Dense)	(None, 9)	1161

=====
 Total params: 828,889
 Trainable params: 828,889
 Non-trainable params: 0
 =====

```
[8]: plot_training_curve(os.path.join(experiment_dest_path, "training_metrics.csv"),
    ↪ ['val_loss', 'val_acc', 'loss', 'acc'])
```



4 Evaluate full precision model

Evaluate performance of the model in full precision to later compare to the quantized one.

```
[ ]: evaluate_model.evaluate(
    model_path=os.path.join(experiment_dest_path, 'model_2d'),
    data=data_path,
    dest_path=experiment_dest_path,
    n_classes=9,
    batch_size=1024,
    noise=[],
    noise_sets=[])
tf.keras.backend.clear_session()
```

```
[10]: show_statistics(os.path.join(experiment_dest_path, "inference_metrics.csv"))
```

```
[10]: accuracy_score balanced_accuracy_score cohen_kappa_score Class_0 \
0      0.842077      0.903612      0.795369  0.843755

      Class_1 Class_2 Class_3 Class_4 Class_5 Class_6 Class_7 \
0  0.783575  0.831801  0.964463  0.99726  0.903327  0.953704  0.856061

      Class_8 inference_time
0  0.998565      7.275701
```

5 Freeze model

Freeze the tensorflow model into the .pb format.

```
[11]: freeze_model.main(model_path=os.path.join(experiment_dest_path, 'model_2d'),
                        output_dir=experiment_dest_path)
```

```
INFO:tensorflow:Froze 61 variables.
INFO:tensorflow:Converted 61 variables to const ops.
Frozen model saved at xillinx_model_compilation_results/experiment_0
```

6 Quantize the model

Perform the quantization by running the `quantize.sh` bash script with appropriate parameters. It executes the `decent_q` command from the Xilinx DNNDK library. The output is the `quantize_eval_model.pb` file and a `deploy_model.pb` file, which should be used for compilation for a specific DPU.

```
[12]: node_names_file = os.path.join(experiment_dest_path,
    ↪ 'freeze_input_output_node_name.json')
frozen_graph_path = os.path.join(experiment_dest_path, 'frozen_graph.pb')
cmd = '../scripts/quantize.sh ' + node_names_file + ' ' \
      + frozen_graph_path + ' ' + data_path + ' ' + \
      '? ,103,1,1' + ' ' + \
```

```

        'ml_intuition.data.input_fn.calibrate_2d_input' + ' ' + \
        '128' + ' ' + experiment_dest_path + \
        ' ' + str(0)
f = open(os.path.join(experiment_dest_path, 'call_output.txt'), 'w')
env = os.environ.copy()
env['PYTHONPATH'] = os.path.dirname(os.getcwd())
subprocess.call(cmd, shell=True, env=env, stderr=f)
f.close()

```

7 Evaluate the quantized model (graph)

Evaluate the performance of the quantized model to check whether there was any loss in performance. Results for the graph are stored in `inference_graph_metrics.csv`.

```

[ ]: graph_path = os.path.join(experiment_dest_path, 'quantize_eval_model.pb')
      evaluate_graph.main(graph_path=graph_path,
                          node_names_path=node_names_file,
                          dataset_path=data_path,
                          batch_size=1024)
      tf.keras.backend.clear_session()

```

```

[14]: show_statistics(os.path.join(experiment_dest_path, "inference_graph_metrics.
      ↪ csv"))

```

```

[14]: accuracy_score balanced_accuracy_score cohen_kappa_score Class_0 \
0      0.874402      0.888489      0.832185 0.908322

      Class_1 Class_2 Class_3 Class_4 Class_5 Class_6 Class_7 \
0 0.89119 0.868037 0.933547 0.998174 0.799121 0.887037 0.712413

      Class_8 inference_time
0 0.998565      21.373184

```

```

[ ]:

```