

model_quantization_xilinx

June 15, 2021

1 Model quantization

If you plan on using this implementation, please cite our work (<https://www.sciencedirect.com/science/article/pii/S0141933119302844>):

@article{Nalepa2020MaM, author = {Jakub Nalepa and Marek Antoniak and Michal Myller and Pablo {Ribalta Lorenzo} and Michal Marcinkiewicz}, title = {Towards resource-frugal deep convolutional neural networks for hyperspectral image segmentation}, journal = {Microprocessors and Microsystems}, volume = {73}, pages = {102994}, year = {2020}, issn = {0141-9331}, doi = {https://doi.org/10.1016/j.micpro.2020.102994}, url = {https://www.sciencedirect.com/science/article/pii/S0141933119302844}}

To perform model quantization, we use the Xilinx DNNDK tool (https://www.xilinx.com/support/documentation/sw_manuals/ai_inference/v1_6/ug1327-dnndk-user-guide.pdf).

```
[1]: import os
import sys
sys.path.append(os.path.dirname(os.getcwd()))
```

```
[4]: import os
import subprocess

import tensorflow as tf

from scripts.quantization import evaluate_graph, freeze_model
from scripts import prepare_data, artifacts_reporter, train_model,
    evaluate_model
from ml_intuition.data.utils import plot_training_curve, show_statistics
```

```
[5]: DEST_PATH = 'xilinx_model_compilation_results'
DATA_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    pavia.npy')
GT_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    pavia_gt.npy')
experiment_dest_path = os.path.join(DEST_PATH, 'experiment_0')
data_path = os.path.join(experiment_dest_path, 'data.h5')
os.makedirs(experiment_dest_path, exist_ok=True)
```

2 Prepare the data

To fit into the the pipeline, the data has to be preprocessed. It is achieved by the `prepare_data.main` function. It accepts a path to a `.npy` file with the original cube as well as the corresponding ground truth. In this example, we randomly extract 250 samples from each class (balanced scenario), use 10% of them as validation set, and extract only spectral information of a pixel. The returned object is a dictionary with three keys: `train`, `test` and `val`. Each of them contains an additional dictionary with `data` and `labels` keys, holding corresponding `numpy.ndarray` objects with the data. For more details about the parameters, refer to the documentation of `prepare_data.main` function (located in `scripts/prepare_data`).

```
[6]: prepare_data.main(data_file_path=DATA_FILE_PATH,
                        ground_truth_path=GT_FILE_PATH,
                        output_path=data_path,
                        train_size=250,
                        val_size=0.1,
                        stratified=True,
                        background_label=0,
                        channels_idx=2,
                        neighborhood_size=None,
                        save_data=True,
                        seed=0)
```

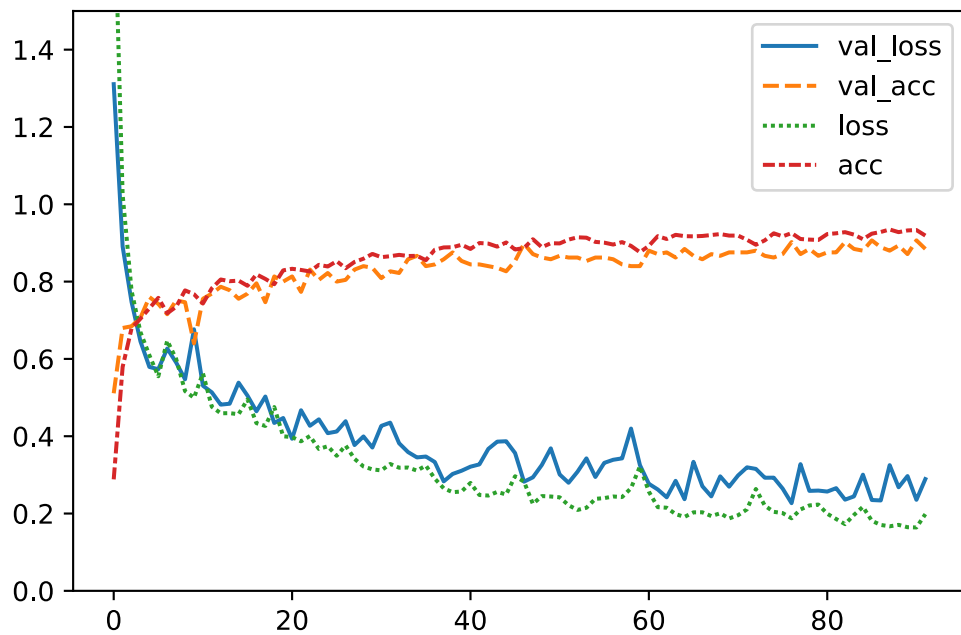
3 Train the model

The function `train_model.train` executed the training procedure. Trained model will be stored under `experiment_dest_path` folder path.

```
[7]: train_model.train(model_name='model_2d',
                        kernel_size=5,
                        n_kernels=200,
                        n_layers=1,
                        dest_path=experiment_dest_path,
                        data=data_path,
                        sample_size=103,
                        n_classes=9,
                        lr=0.001,
                        batch_size=128,
                        epochs=200,
                        verbose=0,
                        shuffle=True,
                        patience=15,
                        noise=[],
                        noise_sets=[])
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 99, 1, 200)	1200
conv2d_1 (Conv2D)	(None, 32, 1, 200)	200200
conv2d_2 (Conv2D)	(None, 14, 1, 200)	200200
conv2d_3 (Conv2D)	(None, 5, 1, 200)	200200
flatten (Flatten)	(None, 1000)	0
dense (Dense)	(None, 200)	200200
dense_1 (Dense)	(None, 128)	25728
dense_2 (Dense)	(None, 9)	1161
Total params: 828,889		
Trainable params: 828,889		
Non-trainable params: 0		

```
[8]: plot_training_curve(os.path.join(experiment_dest_path, "training_metrics.csv"),
    ↪ ['val_loss', 'val_acc', 'loss', 'acc'])
```



4 Evaluate full precision model

Evaluate performance of the model in full precision to later compare to the quantized one.

```
[ ]: evaluate_model.evaluate(  
    model_path=os.path.join(experiment_dest_path, 'model_2d'),  
    data=data_path,  
    dest_path=experiment_dest_path,  
    n_classes=9,  
    batch_size=1024,  
    noise=[],  
    noise_sets=[])  
tf.keras.backend.clear_session()
```

```
[10]: show_statistics(os.path.join(experiment_dest_path, "inference_metrics.csv"))
```

```
[10]: accuracy_score balanced_accuracy_score cohen_kappa_score Class_0 \  
0      0.842077      0.903612      0.795369 0.843755  
  
      Class_1 Class_2 Class_3 Class_4 Class_5 Class_6 Class_7 \  
0 0.783575 0.831801 0.964463 0.99726 0.903327 0.953704 0.856061  
  
      Class_8 inference_time  
0 0.998565      7.275701
```

5 Freeze model

Freeze the tensorflow model into the .pb format.

```
[11]: freeze_model.main(model_path=os.path.join(experiment_dest_path, 'model_2d'),  
                        output_dir=experiment_dest_path)
```

```
INFO:tensorflow:Froze 61 variables.  
INFO:tensorflow:Converted 61 variables to const ops.  
Frozen model saved at xillinx_model_compilation_results/experiment_0
```

6 Quantize the model

Perform the quantization by running the `quantize.sh` bash script with appropriate parameters. It executes the `decent_q` command from the Xilinx DNNDK library. The output is the `quantize_eval_model.pb` file and a `deploy_model.pb` file, which should be used for compilation for a specific DPU.

```
[12]:
```

```

node_names_file = os.path.join(experiment_dest_path,
    ↪ 'freeze_input_output_node_name.json')
frozen_graph_path = os.path.join(experiment_dest_path, 'frozen_graph.pb')
cmd = '../scripts/quantize.sh ' + node_names_file + ' ' \
    + frozen_graph_path + ' ' + data_path + ' ' + \
    '? ,103,1,1' + ' ' + \
    'ml_intuition.data.input_fn.calibrate_2d_input' + ' ' + \
    '128' + ' ' + experiment_dest_path + \
    ' ' + str(0)
f = open(os.path.join(experiment_dest_path, 'call_output.txt'), 'w')
env = os.environ.copy()
env['PYTHONPATH'] = os.path.dirname(os.getcwd())
subprocess.call(cmd, shell=True, env=env, stderr=f)
f.close()

```

7 Evaluate the quantized model (graph)

Evaluate the performance of the quantized model to check whether there was any loss in performance. Results for the graph are stored in `inference_graph_metrics.csv`.

```

[ ]: graph_path = os.path.join(experiment_dest_path, 'quantize_eval_model.pb')
    evaluate_graph.main(graph_path=graph_path,
                        node_names_path=node_names_file,
                        dataset_path=data_path,
                        batch_size=1024)
    tf.keras.backend.clear_session()

```

```

[14]: show_statistics(os.path.join(experiment_dest_path, "inference_graph_metrics.
    ↪ csv"))

```

```

[14]: accuracy_score  balanced_accuracy_score  cohen_kappa_score  Class_0 \
0          0.874402                0.888489                0.832185  0.908322

      Class_1  Class_2  Class_3  Class_4  Class_5  Class_6  Class_7 \
0  0.89119  0.868037  0.933547  0.998174  0.799121  0.887037  0.712413

      Class_8  inference_time
0  0.998565      21.373184

```

```

[ ]:

```