

# model\_noise\_injection\_ensemble

December 2, 2020

## 1 Noise injection into models, ensemble

Inject noise into the model's weights to create an augmented version. There are two scenarios where such operation could be useful: - Inject noise into the model to verify its robustness against unpredicted noise. - Model could be augmented N times, forming an ensemble. Each model makes its own predictions, which are then aggregated to conclude a final prediction.

```
[1]: import os
import sys
sys.path.append(os.path.dirname(os.getcwd()))
```

```
[2]: import os

import tensorflow as tf

from scripts import evaluate_model, prepare_data, artifacts_reporter,
    ↪train_model
```

Specify path to the .npz dataset and ground truth, as well as the output path to store all the artifacts.

```
[3]: DEST_PATH = 'example_results'
DATA_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    ↪pavia.npz')
GT_FILE_PATH = os.path.join(os.path.dirname(os.getcwd()), 'datasets/pavia/
    ↪pavia_gt.npz')
experiment_dest_path = os.path.join(DEST_PATH, 'experiment_0')
os.makedirs(experiment_dest_path, exist_ok=True)
```

## 2 Prepare the data

To fit into the pipeline, the data has to be preprocessed. It is achieved by the `prepare_data.main` function. It accepts a path to a .npz file with the original cube as well as the corresponding ground truth. In this example, we randomly extract 250 samples from each class (balanced scenario), use 10% of them as validation set, and extract only spectral information of a

pixel. The returned object is a dictionary with three keys: `train`, `test` and `val`. Each of them contains an additional dictionary with `data` and `labels` keys, holding corresponding `numpy.ndarray` objects with the data. For more details about the parameters, refer to the documentation of `prepare_data.main` function (located in `scripts/prepare_data`).

```
[4]: data = prepare_data.main(data_file_path=DATA_FILE_PATH,
                             ground_truth_path=GT_FILE_PATH,
                             output_path=None,
                             train_size=250,
                             val_size=0.1,
                             stratified=True,
                             background_label=0,
                             channels_idx=2,
                             neighborhood_size=None,
                             save_data=False,
                             seed=0)
```

### 3 Train the original model

The function `train_model.train` executed the training procedure. Trained model will be stored under `experiment_dest_path` folder path. For details about all arguments, please refer to the documentation of the `train_model.train` function (located in `scripts/train_model`).

```
[5]: train_model.train(model_name='model_2d',
                       kernel_size=5,
                       n_kernels=200,
                       n_layers=1,
                       dest_path=experiment_dest_path,
                       data=data,
                       sample_size=103,
                       n_classes=9,
                       lr=0.001,
                       batch_size=128,
                       epochs=200,
                       verbose=2,
                       shuffle=True,
                       patience=15,
                       noise=[],
                       noise_sets=[])
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 99, 1, 200)	1200
conv2d_1 (Conv2D)	(None, 32, 1, 200)	200200

conv2d_2 (Conv2D)	(None, 14, 1, 200)	200200
conv2d_3 (Conv2D)	(None, 5, 1, 200)	200200
flatten (Flatten)	(None, 1000)	0
dense (Dense)	(None, 200)	200200
dense_1 (Dense)	(None, 128)	25728
dense_2 (Dense)	(None, 9)	1161

=====  
Total params: 828,889

Trainable params: 828,889

Non-trainable params: 0

-----  
Train on 2025 samples, validate on 225 samples

Epoch 1/200

- 2s - loss: 1.8423 - acc: 0.2884 - val\_loss: 1.3107 - val\_acc: 0.5111

Epoch 2/200

- 1s - loss: 1.0250 - acc: 0.5793 - val\_loss: 0.8917 - val\_acc: 0.6800

Epoch 3/200

- 1s - loss: 0.7752 - acc: 0.6790 - val\_loss: 0.7486 - val\_acc: 0.6800

Epoch 4/200

- 1s - loss: 0.6684 - acc: 0.7081 - val\_loss: 0.6481 - val\_acc: 0.7156

Epoch 5/200

- 1s - loss: 0.6096 - acc: 0.7373 - val\_loss: 0.5814 - val\_acc: 0.7689

Epoch 6/200

- 1s - loss: 0.5590 - acc: 0.7565 - val\_loss: 0.6010 - val\_acc: 0.7333

Epoch 7/200

- 1s - loss: 0.6426 - acc: 0.7160 - val\_loss: 0.5996 - val\_acc: 0.7689

Epoch 8/200

- 1s - loss: 0.5781 - acc: 0.7486 - val\_loss: 0.6293 - val\_acc: 0.7200

Epoch 9/200

- 1s - loss: 0.5217 - acc: 0.7699 - val\_loss: 0.5584 - val\_acc: 0.7556

Epoch 10/200

- 1s - loss: 0.4973 - acc: 0.7694 - val\_loss: 0.5960 - val\_acc: 0.6800

Epoch 11/200

- 1s - loss: 0.5231 - acc: 0.7551 - val\_loss: 0.4989 - val\_acc: 0.7822

Epoch 12/200

- 1s - loss: 0.4882 - acc: 0.7921 - val\_loss: 0.6526 - val\_acc: 0.7600

Epoch 13/200

- 1s - loss: 0.4948 - acc: 0.7921 - val\_loss: 0.4600 - val\_acc: 0.7867

Epoch 14/200

- 1s - loss: 0.4596 - acc: 0.8020 - val\_loss: 0.5171 - val\_acc: 0.7556

Epoch 15/200

- 1s - loss: 0.4426 - acc: 0.8020 - val\_loss: 0.4853 - val\_acc: 0.7911

Epoch 16/200  
- 1s - loss: 0.4619 - acc: 0.8059 - val\_loss: 0.5004 - val\_acc: 0.7600

Epoch 17/200  
- 1s - loss: 0.4298 - acc: 0.8247 - val\_loss: 0.4557 - val\_acc: 0.7956

Epoch 18/200  
- 1s - loss: 0.4340 - acc: 0.8064 - val\_loss: 0.4986 - val\_acc: 0.7644

Epoch 19/200  
- 1s - loss: 0.4982 - acc: 0.7837 - val\_loss: 0.4597 - val\_acc: 0.7956

Epoch 20/200  
- 1s - loss: 0.4036 - acc: 0.8291 - val\_loss: 0.4257 - val\_acc: 0.8178

Epoch 21/200  
- 1s - loss: 0.3990 - acc: 0.8356 - val\_loss: 0.4116 - val\_acc: 0.8178

Epoch 22/200  
- 1s - loss: 0.3936 - acc: 0.8291 - val\_loss: 0.5127 - val\_acc: 0.7867

Epoch 23/200  
- 1s - loss: 0.4140 - acc: 0.8277 - val\_loss: 0.4421 - val\_acc: 0.8000

Epoch 24/200  
- 1s - loss: 0.3751 - acc: 0.8351 - val\_loss: 0.4708 - val\_acc: 0.8000

Epoch 25/200  
- 1s - loss: 0.3762 - acc: 0.8395 - val\_loss: 0.4232 - val\_acc: 0.8133

Epoch 26/200  
- 1s - loss: 0.3567 - acc: 0.8459 - val\_loss: 0.4257 - val\_acc: 0.7867

Epoch 27/200  
- 1s - loss: 0.3998 - acc: 0.8286 - val\_loss: 0.4965 - val\_acc: 0.8000

Epoch 28/200  
- 1s - loss: 0.3857 - acc: 0.8316 - val\_loss: 0.4216 - val\_acc: 0.8311

Epoch 29/200  
- 1s - loss: 0.3789 - acc: 0.8380 - val\_loss: 0.4744 - val\_acc: 0.8133

Epoch 30/200  
- 1s - loss: 0.3431 - acc: 0.8538 - val\_loss: 0.4023 - val\_acc: 0.8267

Epoch 31/200  
- 1s - loss: 0.3385 - acc: 0.8563 - val\_loss: 0.3821 - val\_acc: 0.8311

Epoch 32/200  
- 1s - loss: 0.3517 - acc: 0.8474 - val\_loss: 0.4698 - val\_acc: 0.8489

Epoch 33/200  
- 1s - loss: 0.3662 - acc: 0.8365 - val\_loss: 0.3751 - val\_acc: 0.8089

Epoch 34/200  
- 1s - loss: 0.3233 - acc: 0.8573 - val\_loss: 0.3632 - val\_acc: 0.8667

Epoch 35/200  
- 1s - loss: 0.3336 - acc: 0.8607 - val\_loss: 0.3810 - val\_acc: 0.8311

Epoch 36/200  
- 1s - loss: 0.3139 - acc: 0.8696 - val\_loss: 0.3663 - val\_acc: 0.8267

Epoch 37/200  
- 1s - loss: 0.2761 - acc: 0.8904 - val\_loss: 0.3149 - val\_acc: 0.8578

Epoch 38/200  
- 1s - loss: 0.2614 - acc: 0.8854 - val\_loss: 0.3199 - val\_acc: 0.8578

Epoch 39/200  
- 1s - loss: 0.2708 - acc: 0.8859 - val\_loss: 0.3422 - val\_acc: 0.8311

Epoch 40/200  
- 1s - loss: 0.2842 - acc: 0.8815 - val\_loss: 0.3099 - val\_acc: 0.8489  
Epoch 41/200  
- 1s - loss: 0.2766 - acc: 0.8869 - val\_loss: 0.3280 - val\_acc: 0.8444  
Epoch 42/200  
- 1s - loss: 0.2600 - acc: 0.8973 - val\_loss: 0.3292 - val\_acc: 0.8400  
Epoch 43/200  
- 1s - loss: 0.2726 - acc: 0.8770 - val\_loss: 0.3750 - val\_acc: 0.8267  
Epoch 44/200  
- 1s - loss: 0.2696 - acc: 0.8889 - val\_loss: 0.3290 - val\_acc: 0.8667  
Epoch 45/200  
- 1s - loss: 0.2941 - acc: 0.8830 - val\_loss: 0.3248 - val\_acc: 0.8667  
Epoch 46/200  
- 1s - loss: 0.2873 - acc: 0.8844 - val\_loss: 0.3973 - val\_acc: 0.8533  
Epoch 47/200  
- 1s - loss: 0.3035 - acc: 0.8790 - val\_loss: 0.3368 - val\_acc: 0.8533  
Epoch 48/200  
- 1s - loss: 0.2461 - acc: 0.8953 - val\_loss: 0.3060 - val\_acc: 0.8711  
Epoch 49/200  
- 1s - loss: 0.2422 - acc: 0.8919 - val\_loss: 0.3549 - val\_acc: 0.8711  
Epoch 50/200  
- 1s - loss: 0.2303 - acc: 0.9067 - val\_loss: 0.2997 - val\_acc: 0.8622  
Epoch 51/200  
- 1s - loss: 0.2423 - acc: 0.8958 - val\_loss: 0.3090 - val\_acc: 0.8800  
Epoch 52/200  
- 1s - loss: 0.2330 - acc: 0.9027 - val\_loss: 0.2848 - val\_acc: 0.8756  
Epoch 53/200  
- 1s - loss: 0.2117 - acc: 0.9141 - val\_loss: 0.2781 - val\_acc: 0.8578  
Epoch 54/200  
- 1s - loss: 0.2086 - acc: 0.9151 - val\_loss: 0.3018 - val\_acc: 0.8667  
Epoch 55/200  
- 1s - loss: 0.2288 - acc: 0.9086 - val\_loss: 0.2856 - val\_acc: 0.8711  
Epoch 56/200  
- 1s - loss: 0.2319 - acc: 0.9027 - val\_loss: 0.3271 - val\_acc: 0.8578  
Epoch 57/200  
- 1s - loss: 0.2410 - acc: 0.8953 - val\_loss: 0.3243 - val\_acc: 0.8667  
Epoch 58/200  
- 1s - loss: 0.2357 - acc: 0.9042 - val\_loss: 0.4307 - val\_acc: 0.8533  
Epoch 59/200  
- 1s - loss: 0.2814 - acc: 0.8869 - val\_loss: 0.4320 - val\_acc: 0.8533  
Epoch 60/200  
- 1s - loss: 0.3153 - acc: 0.8746 - val\_loss: 0.3284 - val\_acc: 0.8400  
Epoch 61/200  
- 1s - loss: 0.2686 - acc: 0.8894 - val\_loss: 0.2790 - val\_acc: 0.8933  
Epoch 62/200  
- 1s - loss: 0.2201 - acc: 0.9086 - val\_loss: 0.2818 - val\_acc: 0.8889  
Epoch 63/200  
- 1s - loss: 0.2118 - acc: 0.9116 - val\_loss: 0.2571 - val\_acc: 0.8800

```

Epoch 64/200
- 1s - loss: 0.2059 - acc: 0.9160 - val_loss: 0.2581 - val_acc: 0.8711
Epoch 65/200
- 1s - loss: 0.2022 - acc: 0.9160 - val_loss: 0.2651 - val_acc: 0.8756
Epoch 66/200
- 1s - loss: 0.2036 - acc: 0.9195 - val_loss: 0.3508 - val_acc: 0.8444
Epoch 67/200
- 1s - loss: 0.2178 - acc: 0.9126 - val_loss: 0.2374 - val_acc: 0.8889
Epoch 68/200
- 1s - loss: 0.1904 - acc: 0.9284 - val_loss: 0.2619 - val_acc: 0.8844
Epoch 69/200
- 1s - loss: 0.2016 - acc: 0.9185 - val_loss: 0.2581 - val_acc: 0.8978
Epoch 70/200
- 1s - loss: 0.1842 - acc: 0.9269 - val_loss: 0.2634 - val_acc: 0.8622
Epoch 71/200
- 1s - loss: 0.1835 - acc: 0.9269 - val_loss: 0.2828 - val_acc: 0.8889
Epoch 72/200
- 1s - loss: 0.2194 - acc: 0.9116 - val_loss: 0.2562 - val_acc: 0.8711
Epoch 73/200
- 1s - loss: 0.2519 - acc: 0.8983 - val_loss: 0.3043 - val_acc: 0.8489
Epoch 74/200
- 1s - loss: 0.2196 - acc: 0.9101 - val_loss: 0.2614 - val_acc: 0.8800
Epoch 75/200
- 1s - loss: 0.1852 - acc: 0.9294 - val_loss: 0.2702 - val_acc: 0.8756
Epoch 76/200
- 1s - loss: 0.2115 - acc: 0.9072 - val_loss: 0.3366 - val_acc: 0.8533
Epoch 77/200
- 1s - loss: 0.2007 - acc: 0.9269 - val_loss: 0.2438 - val_acc: 0.8933
Epoch 78/200
- 1s - loss: 0.2030 - acc: 0.9146 - val_loss: 0.3068 - val_acc: 0.8800
Epoch 79/200
- 1s - loss: 0.2015 - acc: 0.9121 - val_loss: 0.2778 - val_acc: 0.8800
Epoch 80/200
- 1s - loss: 0.2222 - acc: 0.9096 - val_loss: 0.2870 - val_acc: 0.8578
Epoch 81/200
- 1s - loss: 0.1999 - acc: 0.9269 - val_loss: 0.2449 - val_acc: 0.8889
Epoch 82/200
- 1s - loss: 0.1891 - acc: 0.9235 - val_loss: 0.2787 - val_acc: 0.8889

```

## 4 Evaluate an ensemble of models

To use an ensemble of augmented models, provide a `use_ensemble` argument to `evaluate_model.evaluate` function. Indicate how many copies should be generated with `ensemble_copies` parameter. Lastly, indicate the voting algorithm. It accepts three values:

- **hard** - uses predicted class labels for majority rule voting.
- **soft** - predicts the class label based on the argmax of the sums of the predicted probabilities.

- **classifier** - Use a classifier which accepts predicted probabilities from all the models as features and return the final prediction. The classifier is trained on the train set predictions. Random forest will be used.

Each layer is modified separately. It is achieved by drawing a random number for each of the parameters of the layer. The number is drawn from the normal distribution with provided mean and standard deviation calculated from the layer's parameters, multiplied by 0.1.

```
[ ]: evaluate_model.evaluate(
    model_path=os.path.join(experiment_dest_path, 'model_2d'),
    data=data,
    dest_path=experiment_dest_path,
    n_classes=9,
    batch_size=1024,
    use_ensemble=True,
    ensemble_copies=4,
    voting='hard',
    noise=[],
    noise_sets=[],
    noise_params="{\"mean\": 0, \"std\": None}")
tf.keras.backend.clear_session()
```